

CSCI4180 Tutorial-3

Assignment 1 Review (Part 1)

REN, Yanjing
yjren22@cse.cuhk.edu.hk

Sep. 28, 2022

Outline

- Assignment 1 Overview
 - Specifications
- Hadoop configuration
 - Pseudo distributed mode & Fully distributed mode
- MapReduce Programming
 - WordCount

Assignment 1 Overview

- Four parts
 - Part 1: Getting Started on Hadoop (25%)
 - Part 2: Word Length Count (25%)
 - Part 3: *N*-gram Initial Count (25%)
 - Part 4: Counting *N*-gram Initials Relative Frequencies (25%)

Assignment 1 Overview

- Part 1: Getting Started on Hadoop (25%)
 - Goals
 - Configure Hadoop in *pseudo distributed mode* on single VM
 - **Bouns:** configure Hadoop in *fully distributed mode* with 1 NameNode and 2 DataNodes on pre-assigned 3 VMs
 - Run WordCount program
 - Demo
 - Show your VM and Hadoop configurations
 - Run WordCount with 2 datasets: **KJV Bible** and **the complete works of Shakespeare**

Assignment 1 Overview

- Part 2: Word Length Count (25%)
 - Goals
 - Extend **WordCount** program to **WordLengthCount**
 - Implement **in-mapper combining** to optimize your program
 - Program output: **(length, count)**, separated by an **empty space**.
 - **Important Notes**
 - Double check your file name (**WordLengthCount.java**)
 - Do **NOT** specify the package (e.g., “package <pkg_name>”) in your program

Assignment 1 Overview

- Part 3: *N*-gram Initial Count (25%)
 - Goals
 - Extend ***WordLengthCount*** program to ***NgramInitialCount***
 - Input
 - Additional Parameters
 - N: number of words in N-gram
 - Data
 - You may assume that the datasets contain only printable ASCII characters

Assignment 1 Overview

- Part 3: *N*-gram Initial Count (25%)
 - Definitions
 - **Word**
 - A sequence of English alphabets characters (Case-sensitive)
 - **Ngram**
 - A sequence of **N consecutive words**
 - The word in the end of a line and the word in the beginning of the next line also form an N-gram
 - **Ngram Initial**
 - The first letters of the N words of the N-gram

Assignment 1 Overview

- Part 3: *N*-gram Initial Count (25%)
 - **Delimiter**
 - Non-alphabet characters
 - {ASCII } \ {alphabetic characters}
 - Multiple consequent non-alphabet characters are treated as a single delimiter

Assignment 1 Overview

- Part 4: Counting N -gram Initials Relative Frequencies (25%)
 - Extend ***NgramInitialCount*** program to ***NgramInitialRF***
 - Relative frequency = $\text{count}(\text{"X Y Z"}) / \text{count}(\text{"X *"})$
 - "*" stands for the remaining $N - 1$ characters of **any** N -gram Initials.
 - In this example, $N = 3$
 - Both "X Y Z" and "X *" are 3-gram words
 - Output
 - **ONLY** output sequences with frequency $\geq \theta$

Outline

- Assignment 1 Overview
 - Specifications
- Hadoop configuration
 - Pseudo distributed mode & Fully distributed mode
- MapReduce Programming
 - WordCount

Hadoop Configuration

- core-site.xml

pseudo distributed mode

```
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://localhost:9000</value>
</property>
```

fully distributed mode

```
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://vm1:9001</value>
</property>
```

Localhost -> vm1

Follow the [link](#) to set up pseudo distributed mode

We also provide a set of configuration files on Blackboard: [Link](#)

Hadoop Configuration

- `hdfs-site.xml`

pseudo distributed mode

```
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
```

fully distributed mode

```
<property>
<name>dfs.replication</name>
<value>2</value>
</property>
```

Replication 1 -> 2

- `slaves`

localhost-> vm2 vm3

Outline

- Assignment 1 Overview
 - Specifications
- Hadoop configuration
 - Pseudo distributed mode & Fully distributed mode
- MapReduce Programming
 - WordCount

MapReduce Programming

- Basic Structure

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

Depends on the type of your Mapper and Reducer's input and output

MapReduce Programming

- Basic Structure

- public class **WordCount**

```
public static class TokenizerMapper  
    extends Mapper<Object, Text, Text, IntWritable>
```

The **map** function is put inside this class

```
public static class IntSumReducer  
    extends Reducer<Text, IntWritable, Text, IntWritable>
```

The **reduce** function is put inside this class

```
public static void main(String[] args) throws Exception
```

The **starting** point of the job, which includes job configuration

MapReduce Programming

- Mapper

- Refer to [Mapper API](#)

Method and Description

cleanup `org.apache.hadoop.mapreduce.Mapper.Context context`)

Called once at the end of the task.

map **KEYIN** key, **VALUEIN** value, `org.apache.hadoop.mapreduce.Mapper.Context context`)

Called once for each key/value pair in the input split.

run `(org.apache.hadoop.mapreduce.Mapper.Context context)`

Expert users can override this method for more complete control over the execution of the Mapper.

setup `org.apache.hadoop.mapreduce.Mapper.Context context`)

Called once at the beginning of the task.

MapReduce Programming

- Mapper

- Class Mapper<KEYIN,VALUEIN,KEYOUT,VALUEOUT>
 - We use **Generics** to specify the data type of the input key-value pair (KEYIN, VALUEIN) and output key-value pair (KEYOUT, VALUEOUT)
 - We can use the following types [here](#)
 - By default, the KEYIN and VALUEIN of the mapper class is **LongWritable** and **Text** respectively, which represents *the line number* and *the whole line* in the input

MapReduce Programming

- Reducer

- Refer to [Reducer API](#)

| Method and Description |
|---|
| cleanup (org.apache.hadoop.mapreduce.Reducer.Context context) Called once at the end of the task. |
| reduce (KEYIN key, Iterable<VALUEIN> values, org.apache.hadoop.mapreduce.Reducer.Context context) This method is called once for each key. |
| run (org.apache.hadoop.mapreduce.Reducer.Context context) Advanced application writers can use the run(org.apache.hadoop.mapreduce.Reducer.Context) method to control how the reduce task works. |
| setup (org.apache.hadoop.mapreduce.Reducer.Context context) Called once at the start of the task. |

MapReduce Programming

- Reducer

- Class Reducer<KEYIN,VALUEIN,KEYOUT,VALUEOUT>
 - Like mapper, we need to specify the data type of the input key-value pair (KEYIN, VALUEIN) and output key-value pair (KEYOUT, VALUEOUT)
 - The mapper's (KEYOUT, VALUEOUT) should be the same as the reducer's (KEYIN, VALUEIN)
 - This is because the input of reducers comes from the output of mappers

MapReduce Programming

- Configure Job

```
Job job = Job.getInstance(conf, "word count");  
job.setJarByClass(WordCount.class);  
job.setMapperClass(TokenizerMapper.class);  
job.setCombinerClass(IntSumReducer.class);  
job.setReducerClass(IntSumReducer.class);  
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);  
FileInputFormat.addInputPath(job, new Path(args[0]));  
FileOutputFormat.setOutputPath(job, new Path(args[1]));  
System.exit(job.waitForCompletion(true) ? 0 : 1);
```

Name of the Job

MapReduce Programming

- Configure Job

```
Job job = Job.getInstance(conf, "word_count");
job.setJarByClass(WordCount.class);
job.setMapperClass(TokenizerMapper.class);
job.setCombinerClass(IntSumReducer.class);
job.setReducerClass(IntSumReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
```

Should be the same as the name of public class

MapReduce Programming

- Configure Job

```
Job job = Job.getInstance(conf, "word count");
job.setJarByClass(WordCount.class);
job.setMapperClass(TokenizerMapper.class);
job.setCombinerClass(IntSumReducer.class);
job.setReducerClass(IntSumReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
```

The class name of Mapper, Combiner and Reducer

MapReduce Programming

- Configure Job

```
Job job = Job.getInstance(conf, "word count");
job.setJarByClass(WordCount.class);
job.setMapperClass(TokenizerMapper.class);
job.setCombinerClass(IntSumReducer.class);
job.setReducerClass(IntSumReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
```

Specify the data types of the output key-value pair. They should match with that of the output key-value pair of Reducer.

MapReduce Programming

- Configure Job

```
Job job = Job.getInstance(conf, "word count");
job.setJarByClass(WordCount.class);
job.setMapperClass(TokenizerMapper.class);
job.setCombinerClass(IntSumReducer.class);
job.setReducerClass(IntSumReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
```

Specify the input and output directory.
Here, we use command-line arguments
to define the path.

MapReduce Programming

- Configure Job

```
Job job = Job.getInstance(conf, "word count");
job.setJarByClass(WordCount.class);
job.setMapperClass(TokenizerMapper.class);
job.setCombinerClass(IntSumReducer.class);
job.setReducerClass(IntSumReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
```

Submit the job and wait for it to complete

MapReduce Programming

- **Notes**

- If you write codes in VMs, please **backup** your code
- Check the assignment specifications very carefully, especially the **file names, input / output formats**
 - Make sure we can compile and run your codes, and double check the format of your program output

Q & A