

CSCI4180: Tutorial 6

Assignment 2 Review (Part 1)

Huancheng Puyang

Department of Computer Science & Engineering
The Chinese University of Hong Kong

2022.10.26

Outline

➤ Problem

- Single-source shortest path length by parallel Dijkstra's algorithm

➤ Three main modules

- PDPreProcess.java
- PDNodeWritable.java
- ParallelDijkstra.java

➤ Implementation hints

- Emit a node
- Stop condition
- Chain MapReduce jobs

Outline

➤ Problem

- Single-source shortest path length by parallel Dijkstra's algorithm

➤ Three main modules

- PDPreProcess.java
- PDNodeWritable.java
- ParallelDijkstra.java

➤ Implementation hints

- Emit a node
- Stop condition
- Chain MapReduce jobs

Problem

➤ Single-source shortest path

- Find shortest paths from a source node to all other nodes in the graph
- Formalization
 - Given a graph $G = (V, E)$ and a source node $v_s \in V$
 - Find the shortest path length (distance) from v_s to every other reachable node in V
- Breadth-first-search (BFS): smallest hop counts (unweighted)
- Dijkstra's algorithm: lowest weights (weighted)

Problem (Cont.)

➤ Parallel Dijkstra's algorithm

- Refer to Slides 35-49 of Lecture 5 for details
- Similar workflow as parallel BFS, yet with some differences
 - Consider general weights (weighted) instead of unweighted.
 - Consider stop iteration: how to know that all distances have been found?
 - Maintain a counter inside MapReduce program for some statistic, e.g., the number of nodes to be processed after each iteration (detailed later)

Problem (Cont.)

➤ Dataset: Twitter social graph

- URL: <https://anlab-kaist.github.io/traces/WWW2010>
- Original per-line structure, e.g., user1 user2
 - It means the relationship that user1 follows user2
- Modification for part 2 of assignment 2
 - User → node; relationship → edge
 - Each user is represented as a node
 - Each node takes a unique positive integer as the **node ID**
 - E.g., for user1 follows user2, an edge is created from node1 to node2 in the graph
 - Original dataset is unweighted, we randomly assign a **positive integer** between 1 and 50 inclusively as the weight for each edge, e.g., user1 user2 → node1 node2 30
 - Original dataset is tens of GB in total, we build a smaller one by sampling

Problem (Cont.)

- Dataset: Twitter social graph (cont.)
 - We have uploaded the dataset to Blackboard
 - Each part has a small case and a large case
 - Small case: example given by specification (tens of bytes)
 - Large case: sampled from Twitter dataset (tens of MB)
 - You can use the dataset for debug and test

Outline

➤ Problem

- Single-source shortest path length by parallel Dijkstra's algorithm

➤ Three main modules

- PDPreProcess.java
- PDNodeWritable.java
- ParallelDijkstra.java

➤ Implementation hints

- Emit a node
- Stop condition
- Chain MapReduce jobs

Three Main Modules

➤ PDPreProcess.java for parsing input

- Example for parsing input
 - Model following lines to $G = (V, E)$
 - 1173 1173 10
 - Add node 1173 to V
 - No edge is added to E
 - 1173 6267522 14
 - Add node 1173 to V
 - Add node 6267522 to V
 - Add edge from 1173 to 6267522 to E , whose weight is 14



Three Main Modules (Cont.)

➤ PDPreProcess.java for parsing input

- Example for parsing input
 - Input example from specification:

1	2	7
1	3	20
2	3	3
3	1	5
4	1	9
5	6	10

Input



Three Main Modules (Cont.)

➤ PDPPreProcess.java for parsing input

- **Input** format → **adjacency list** format
- A separate Mapper/Reducer to do the transformation
- For the input example from specification:

1	2	7
1	3	20
2	3	3
3	1	5
4	1	9
5	6	10

Input



1	(2 7), (3 20)
2	(3 3)
3	(1 5)
4	(1 9)
5	(6 10)

Adjacency List

Three Main Modules (Cont.)

➤ PDNodeWritable.java for node structure

- Take a node v_k as an example, where v_s is the source node
 - Stores the node ID v_k to identify which node the structure belongs to
 - Stores the distance (shortest path length) from v_s to v_k
 - Keeps track of the neighbors of v_k by adjacency list
 - Keeps track of the ID of previous node leading to the minimum distance from v_s to v_k
- Hints
 - Carefully design your structure
 - Proper constructor
 - Proper type for node ID, distance and adjacency list
 - Feel free to design your own format for right functionality

Three Main Modules (Cont.)

➤ ParallelDijkstra.java for main module

- Mapper
 - For each node v_k
 - Receive node structure
 - Distance from v_s to v_k and adjacency list of v_k including distances from v_k to its neighbors
 - Calculate the distances from v_s to v_k 's neighbors through v_k
 - Emit the node structure of v_k itself
 - Emit the distances from v_s to v_k 's neighbors through v_k
 - Consider how to design your implementation

Three Main Modules (Cont.)

➤ ParallelDijkstra.java for main module

- Reducer
 - For each node v_k
 - Receive the node structure of v_k and the distances from v_s to v_k through the predecessors
 - Find the smallest value and the corresponding previous node
 - Update the distance of v_k in node structure using the smallest value
 - Record the previous node in node structure
 - Think about how to fix the conflicting information
 - **NOTE:** Emitted node can be used by the Mappers in the next iteration

Outline

➤ Problem

- Single-source shortest path length by parallel Dijkstra's algorithm

➤ Three main modules

- PDPreProcess.java
- PDNodeWritable.java
- ParallelDijkstra.java

➤ Implementation hints

- Emit a node
- Stop condition
- Chain MapReduce jobs

Implementation Hints

➤ Emit a node

- The Java Way

- Implements a toString() method and a fromString() method in PDNodeWritable class
 - You can convert a Node to String, and parse a Node from String
- Emit node.toString() simply by a [Text](#) and parse a Node by fromString() from a [Text](#)
 - See the hyperlink for how to convert between Text and String
- NOTE: functions for parsing String like String.split() could be slow

- The Hadoop Way (Recommended)

- Implement [Writable](#) in PDNodeWritable class
 - You need to provide implementation for methods, e.g., readFields, write, and read
- Emit a Node as the value directly
- Could be faster than parsing String in the Java way, worth trying if you want to go for Bonus!

Implementation Hints (Cont.)

➤ Stop condition

- Command line argument: iterations
 - Indicate the number of MapReduce iterations
 - Positive integer
- Stop condition:
 - Iterations > 0 → Stop after *iterations* runs (maximum number) of MapReduce (maybe earlier if all nodes are processed)
 - You can use [Counters](#) to decide if all reachable nodes have been processed (Slide 4)
- Count the number of iterations
 - You can use [Counters](#) in Hadoop

Implementation Hints (Cont.)

➤ Stop condition (cont.)

- Basic operations about Hadoop Counter
 - Declare Counter
 - `public static enum ReachCounter { COUNT };`
 - Increment Counter
 - `context.getCounter(ReachCounter.COUNT).increment(1);`
 - Retrieve Counter Value
 - `long reachCount =
job.getCounters().findCounter(ParallelDijkstra.ReachCounter.COUNT).getValue();`

Implementation Hints (Cont.)

➤ Chain MapReduce jobs

- Each “Map + Reduce” only explores one step further from v_s , we need to repeat this process
 - Map1, Reduce1, Map2, Reduce2, Map3...
- How to chain MapReduce jobs?

Implementation Hints (Cont.)

➤ Chain MapReduce jobs (cont.)

- The Java Way
 - Drive jobs in the main() function using loops, counter and conditions
 - Different **Configuration** and **Job Object** in each iteration
 - Set job.waitForCompletion(true);
 - Jobs communicate by writing and reading intermediate files on HDFS
 - For Job_i :
 - FileInputFormat.addInputPath(*OutputPath of Job_{i-1}*)

Implementation Hints (Cont.)

➤ Chain MapReduce jobs (cont.)

- The Hadoop Way
 - Use [JobControl](#)
 - Add jobs to JobControl
 - `JobControl jc = new JobControl();`
 - `jc.addJob(job1);`
 - `jc.addJob(job2);`
 - Add dependency (e.g. job2 depends on job1)
 - `job2.addDependingJob(job1)`
 - Run JobControl
 - `jc.run()`
 - NOTE: YARN will schedule those jobs for you wisely

Implementation Hints (Cont.)

➤ Remarks

- Cleanup
 - Transform the output of the last iterator to the required format
 - Only outputs the tuple for the nodes which are reachable from the source node
 - Feel free to use another set of Mapper/[Reducer] to do this step
- Tips
 - Test your program correctness with hand-craft test cases (loops, directed edge, etc.)
 - **NOTE:** Please start early!

**Thank You
Q&A**