

CSCI4180:Tutorial-7

Assignment 2 Review (Part 2)

Huancheng Puyang

Department of Computer Science & Engineering
The Chinese University of Hong Kong

2022.11.02

Outline

- Problem
 - Measure web page quality by PageRank algorithm
- Three main modules
 - PRPreProcess.java
 - PRNodeWritable.java
 - PageRank.java
- Implementation hints
- Submission

Outline

➤ Problem

- Measure web page quality by PageRank algorithm

➤ Three main modules

- PRPreProcess.java
- PRNodeWritable.java
- PageRank.java

➤ Implementation hints

➤ Submission

Problem

➤ Model of web pages

- **Network of pages** are modeled as a **directed graph**
 - **Page** → **node**
 - Each page is represented as a node
 - Each node takes a unique positive integer as the node ID
 - **Hyperlink** → **edge**
 - When there is a hyperlink in page1 that leads user to page2, there is an edge from node1 to node2 in directed graph

➤ Dataset

- A small one with a few lines and a large one sampled from Twitter dataset
- Get it from Blackboard and use it for debug and test

Problem (Cont.)

➤ Example

Edge (node 1, node 2)

⚠ 不安全 cse.cuhk.edu.hk/~pclee/csci4180/

Week 2 9/14, 9/16	Lecture 2: Overview of Ha Required readings: <ul style="list-style-type: none">• LD Ch. 2 Optional readings: <ul style="list-style-type: none">• White Ch. 1-4• Apache Hadoop 2.7.3• Ghemawat et al., "The Go• Dean and Ghemawat, "Ma• Vavilapalli et al., "Apache
----------------------	--

Node 1

← → ↻ 🏠 🔒 hadoop.apache.org/docs/r2.7.3/



Apache > Hadoop

→ **General**

Overview

- Single Node Setup
- Cluster Setup
- Hadoop Commands
- Reference
- FileSystem Shell

Apache Hadoop

Apache Hadoop 2.7.3 is a minor release

Here is a short overview of the major fe

Node 2

Problem (Cont.)

➤ PageRank algorithm

- Measure the quality of web pages by **iterative** computation

➤ Simple sketch of algorithm (Slide 55 of lecture 5)

- Caution: no dangling nodes and no random jump factor
- **Page_i** starts with seed **PR_i** values (e.g., each page has equal PageRank, i.e., $PR_i = \frac{1}{|G|}$)
- **Page_i** distributes **PR_i** “credit” to all pages it links to
 - **Page_i** also receives credits from its predecessors by multiple in-bound links
- **Page_i** adds up “credit” from multiple in-bound links to compute **PR_i'**
- Iterate until values converge

Problem (Cont.)

➤ Overview of part 3

- Input format
 - Each line: <node ID 1> <node ID 2> <weight>
 - Similar as part 1, yet we **ignore weight** in part 2
 - Example: “5 2 14” means an edge from node 5 to node 2 **without weight**
- Command line arguments
 - *iteration*: number of iterations for main loop (without any other stop condition)
 - *threshold*: minimal PageRank value for a node to output
 - *infile*: path of input file
 - *outdir*: path of output

Problem (Cont.)

➤ Overview of part 3 (Cont.)

- Output format
 - Each line: <node ID> <PageRank value>
 - NOTE: you need to output all nodes whose PageRank value is above threshold

Problem (Cont.)

➤ Overview of part 3 (Cont.)

- Pre-processing (PRPreProcess.java)
- PageRank algorithm (PageRank.java)
 - Map: distribute the “credits” to all pages that the current page links to
 - Reduce: sum up the “credits” from all in-bound links
- Output
 - Transform the results into required output format
- Similar as workflow of part 2
 - Pre-processing (PDPreProcess.java)
 - ParallelDijkstra algorithm (ParallelDijkstra.java)
 - Map: emit the distances going through the current node
 - Reduce: update shortest distance according to received values
 - Output

Outline

- Problem
 - Measure web page quality by PageRank algorithm
- Three main modules
 - PRPreProcess.java
 - PRNodeWritable.java
 - PageRank.java
- Implementation hints
- Submission

Three Main Modules

- PRPreProcess.java for parsing input
 - Parse input file into graph $G = (V, E)$
 - Represented by adjacency list
 - Similar as part 2 (see tutorial 6)
 - Simple example


4	1	1
1	2	1
1	3	1
2	3	1
3	1	1

Input



1	(2 3)
2	(3)
3	(1)
4	(1)

Adjacency list



Three Main Modules (Cont.)

- PRNodeWritable.java for node structure
 - For each node
 - Variable
 - Node ID
 - Adjacency list
 - PageRank value
 - ...
 - Method
 - toString() and fromString()
 - readFields() and write()
 - Similar as PDNodeWritable.java in part 2 (see tutorial 6)
 - Both include node ID, adjacency list, and conversion methods
 - Difference: part 2 needs to store shortest distance and previous node

Three Main Modules (Cont.)


- PageRank.java for main module
 - Launch a MapReduce job for pre-processing module
 - **Main loop** to calculate per-node PageRank value iteratively
 - Cleanup for requirement output format
 - Similar as ParallelDijkstra.java in part 2
 - Launch a MapReduce job for pre-processing module
 - Main loop to update shortest distance by Dijkstra's algorithm
 - Cleanup for requirement output format

Three Main Modules (Cont.)

- Main loop in PageRank.java (for each iteration)
 - MapReduce job to update PageRank value without random jump
 - Mapper
 - Get per-node PageRank value as input
 - For 1st iteration: initialize per-node PageRank value as $1/N$, where N is # of nodes
 - Q: How to get the total number of nodes?
 - You can use [Counters](#) to count the number of nodes in pre-processing module

Three Main Modules (Cont.)

- Main loop in PageRank.java (for each iteration)
 - MapReduce job to update PageRank value without random jump
 - Mapper
 - Get per-node PageRank value as input
 - For 1st iteration: initialize per-node PageRank value as $1/N$, where N is # of nodes



Three Main Modules (Cont.)

- Main loop in PageRank.java (for each iteration)
 - MapReduce job to update PageRank value without random jump
 - Mapper
 - Get per-node PageRank value as input
 - For 1st iteration: initialize per-node PageRank value as $1/N$, where N is # of nodes
 - For other iterations: use results of previous iteration as the input of current iteration

Three Main Modules (Cont.)


- Main loop in PageRank.java (for each iteration)
 - MapReduce job to update PageRank value without random jump
 - Mapper
 - Get per-node PageRank value as input
 - Evenly distribute PageRank value of each node to its successors

```
1: class MAPPER
2:   method MAP(nid n, node N) Evenly divide PageRank value of node n
3:      $p \leftarrow N.PAGERANK / |N.ADJACENCYLIST|$ 
4:     EMIT(nid n, N) Emit node structure of node n ▷ Pass along graph structure
5:     for all nodeid m ∈ N.ADJACENCYLIST do
6:       EMIT(nid m, p) ▷ Pass PageRank mass to neighbors
```

Emit divided PageRank value to each successor

Three Main Modules (Cont.)

- Main loop in PageRank.java (for each iteration)
 - MapReduce job to update PageRank value without random jump
 - Mapper
 - Get per-node PageRank value as input
 - Evenly distribute PageRank value of each node to its successors




Three Main Modules (Cont.)

- Main loop in PageRank.java (for each iteration)
 - MapReduce job to update PageRank value without random jump
 - Reducer: recover node structure and update PageRank value
 - **NOTE:** the value type of reducer is node structure (PRNodeWritable.java)
 - Provide not only **full information** of node m , but also **PageRank value** given by each predecessor → **carefully design!** (also in part 1: not only **information** but also **distances**)

```
1: class REDUCER
2:   method REDUCE(nid  $m$ , [ $p_1, p_2, \dots$ ])
3:      $M \leftarrow \emptyset$ 
4:     for all  $p \in \text{counts } [p_1, p_2, \dots]$  do
5:       if IsNode( $p$ ) then Node structure of node  $m$ 
6:          $M \leftarrow p$  ▷ Recover graph structure
7:       else
8:          $s \leftarrow s + p$  ▷ Sum incoming PageRank contributions
9:          $M.\text{PAGE RANK} \leftarrow s$  New PageRank value of node  $m$  got from predecessors
10:    EMIT(nid  $m$ , node  $M$ )
```

Three Main Modules (Cont.)

- Main loop in PageRank.java (for each iteration)
 - MapReduce job to update PageRank value without random jump
 - Reducer: recover node structure and update PageRank value



Three Main Modules (Cont.)

- Main loop in PageRank.java (for each iteration)
 - We do not consider the dangling nodes in this assignment.
 - There is no dangling node in the large dataset.

Outline

- Problem
 - Measure web page quality by PageRank algorithm
- Three main modules
 - PRPreProcess.java
 - PRNodeWritable.java
 - PageRank.java
- **Implementation hints**
- Submission

Implementation Hints

➤ Datatype of PageRank value

- For example, for a graph of six nodes, each node will have initial PageRank value of $1/6$ (0.166...)
- For accuracy, you should use **double** when calculating PageRank value

➤ Stop condition

- Stop after a given number of iterations

➤ Cleanup

- Remember to transform your result to required output format at last, otherwise you might **lose** grade

Implementation Hints (Cont.)

- How to chain multiple iterations?
 - The output in iteration(k) serves as the input in iteration(k+1)
- How to obtain the total number of nodes?
 - See Counter [API](#) and [Example](#)
- How to pass parameters for different jobs?
 - Set parameter in job configuration

Outline

- Problem
 - Measure web page quality by PageRank algorithm
- Three main modules
 - PRPreProcess.java
 - PRNodeWritable.java
 - PageRank.java
- Implementation hints
- **Submission**

Submission

- Submit at least the following files (Additional files allowed)
 - Part 2
 - ParallelDijkstra.java
 - PDNodeWritable.java
 - PDPreProcess.java
 - Part 3
 - PageRank.java
 - PRNodeWritable.java
 - PRPreProcess.java
 - **Group declaration form**
 - [http://www.cuhk.edu.hk/policy/academichonesty/Eng_hm_files_\(2013-14\)/p10.htm](http://www.cuhk.edu.hk/policy/academichonesty/Eng_hm_files_(2013-14)/p10.htm)

Submission (Cont.)

- Submit at least the following files
 - Follow the submission instructions under the [Assignments section](#).
 - You can enter your source code directory, and get the tarball by
 - `tar cvzf asgn2-{SID}.tar.gz *.java {your_declaration_form}`
 - **Double-check** your tarball to ensure that it includes all required files
 - Copy `asgn2-*SID*.tar.gz` to a temporary directory
 - Extract it for double-check: `tar xvzf asgn2-{SID}.tar.gz`

**Thank You
Q&A**