

Lecture 7: Key-Value Stores (Part 2)

CSCI4180

Patrick P. C. Lee

Consistency is Hard to Achieve

Initially, $x = 0$; $y = 0$

Process A

```
PUT("x", 1)
PUT("y", 1)
print GET("x")
```

Process B

```
while (GET("y") != 1) {
    sleep()
}
print GET("x")
```

- What are the outputs of Process A and Process B?
- In a single machine, we can use **memory barrier** to solve the problem
- What about in a distributed system? Any difficulty to put a “memory barrier”?

Outline

- NoSQL Origin
 - CAP Algorithm
- Case study: Amazon Dynamo

CAP Theorem

- CAP Conjecture by Eric Brewer in 2000
 - **C: consistency**, by which a shared and replicated data item appears as a single, up-to-date copy
 - **A: availability**, by which updates will always be eventually executed
 - **P: partition tolerance**, by which partitioning of nodes can be tolerated
 - Any networked system providing shared data can provide only two of the three properties
- It is later proven to be correct by Gilbert and Lynch in 2002
 - In a network subject to communication failures, it is impossible to realize an atomic read/write shared memory that guarantees a response to every request [Gilbert and Lynch, 2002]

CAP Theorem

➤ **C & P**: e.g., distributed database systems

- Consistency?
 - When a user writes data, the same piece of data is written to all storage servers
- Partition tolerance?
 - All storage nodes are guaranteed to have the same piece of data, even though part of the system is disconnected
- Availability?
 - Cannot hold! Why?

CAP Theorem

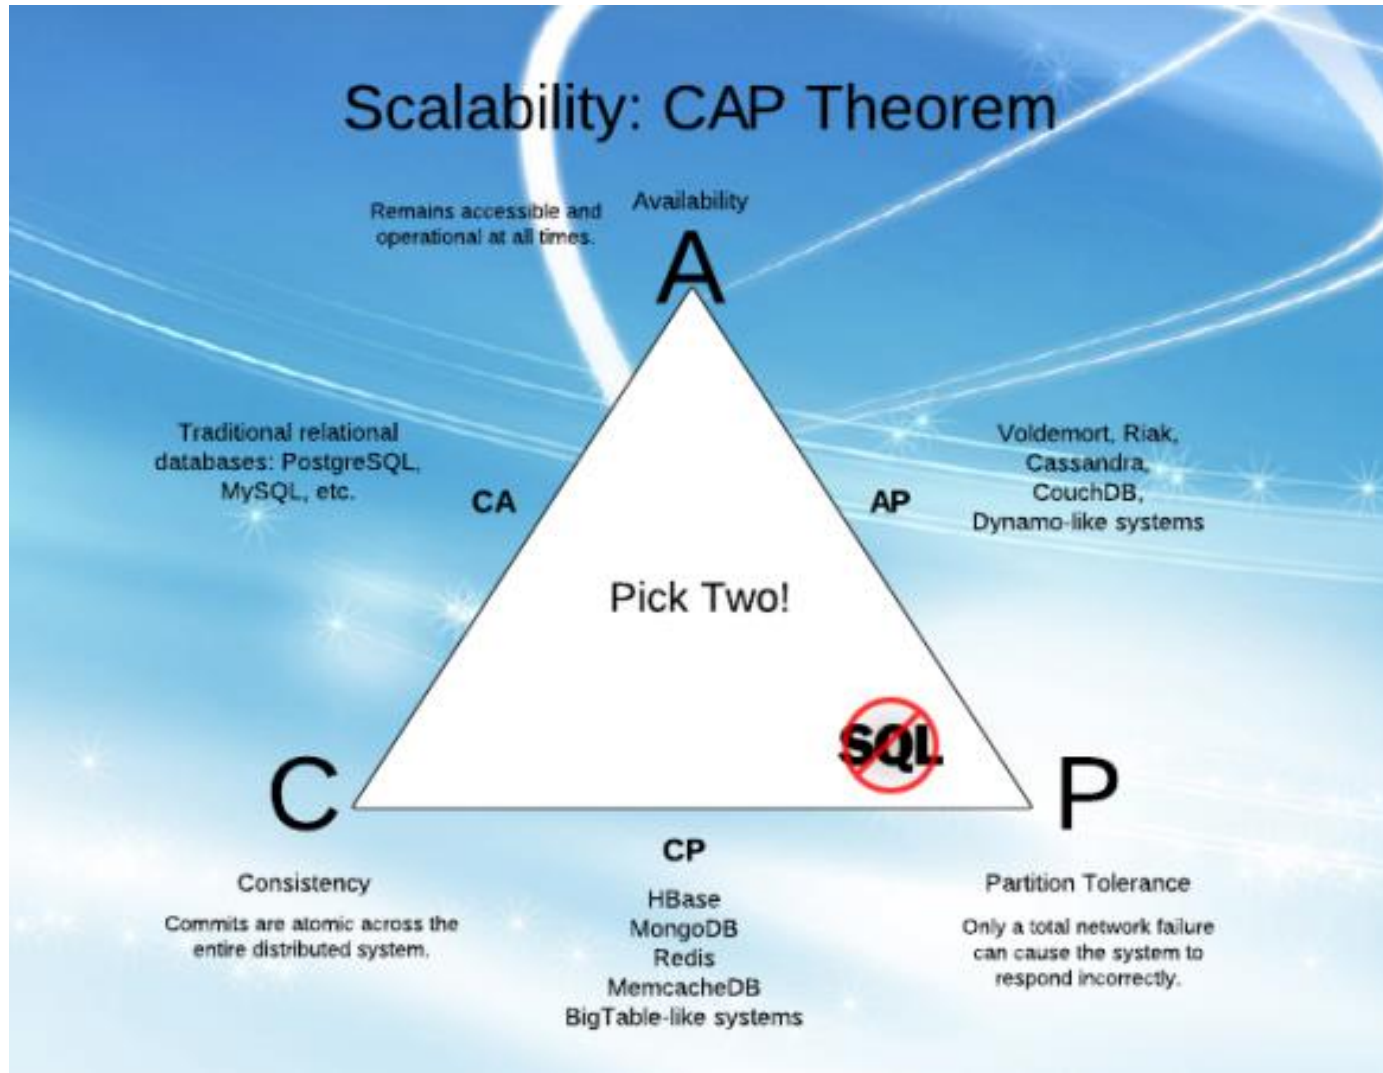
➤ **A & P**: e.g., DNS system

- Availability?
 - YES! Every node in system will always be available no matter the <IP-name> pairs is valid or not
 - e.g., when using “**no-ip.org**”, it is always available, although the results may be invalid
 - Note that **invalid != inconsistent!**
- Partition tolerance?
 - By design, every user (or machine) has its own DNS server on its own network, so network partitioning is not a problem.
- Consistency?
 - Cannot hold. Why?

CAP Theorem

- **C & A**: e.g., single-node database
 - Consistency?
 - Availability?
 - Partition tolerance?
 - Cannot hold. Why?

CAP Theorem



CAP in Reality

- “2 of 3” is misleading
 - It’s really between C and A only
- Reality:
 - Partitions are rare; no need to forfeit C or A
 - Partitions manifest themselves as high delays
 - Shouldn’t just give up either C or A for partition tolerance
 - Correct implementation: detect partitions, enter degraded mode to limit operations, initiate recovery when partitions are fixed
- Choice between C and A varies across operations and data → how to trade between C and A is critical

Outline

- NoSQL Origin
 - CAP Algorithm
- Case study: Amazon Dynamo

Amazon Dynamo

- Dynamo is designed as a **highly available, key-value storage** system
 - Used to power parts of Amazon storage services (e.g., S3)
- In an infrastructure with millions of components, something is always failing!
 - Failure is the normal case
- Good news: many services on Amazon's platform only need **primary-key** access
 - RDBMS is unnecessary

Assumptions and Requirements

- Simple read/write operations, uniquely identified by **keys**
 - Objects tend to be small ($< 1\text{MB}$)
- ACID gives poor availability
 - *Atomicity, Consistency, Isolation, Durability*
 - Use weaker consistency (C) for higher availability
- At least 99.9% of read/write operations must be performed within a few hundred milliseconds:
 - Avoid routing requests through multiple nodes

Assumptions and Requirements

- No security issues
 - Under single authority
- Applications can configure Dynamo for desired latency & throughput, balancing performance, cost, availability, durability guarantees
- Dynamo aims to be **always writable**
 - Rejecting updates is bad in customer experience (e.g., in shopping cart updates)
 - Push complexity of conflict resolution to reads

Design Principles

- Incremental scalability
 - System should be able to grow by adding a storage host (node) at a time
- Symmetry
 - Every node has the same set of responsibilities
- Decentralization
 - Favor decentralized techniques over central coordinators
- Heterogeneity
 - Workload partitioning should be proportional to capabilities of servers

Compared to BigTable

- Dynamo targets apps that only need key/value access with a primary focus on high availability
 - key-value store versus column-store (column families and columns within them)
 - Bigtable: distributed DB built on GFS
 - Dynamo: distributed hash table
 - Updates are not rejected even during network partitions or server failures

Consistency vs. Availability

- Strong consistency and high availability cannot be achieved simultaneously with high performance
- Replication of objects:
 - Good:
 - High availability
 - High read performance
 - Bad:
 - Need to update all of them to maintain consistency
- Dynamo enforces **eventual consistency**
 - Propagate updates to all replicas “eventually”
 - Trade weak consistency for availability and performance

Consistency and Availability

- In eventual consistency model, reads may return old data that is yet updated
- Who resolves conflicts?
 - Choices: the data store or the application
 - Data store
 - application-unaware, so choices limited simple policy, such as "last write wins"
 - Application
 - app is aware of the meaning of the data
 - can do application-aware conflict resolution
 - e.g., merge shopping cart versions to get a unified shopping cart.
 - fall back on "last write wins" if app doesn't want to bother

Dynamo's Design

Problem	Technique	Advantage
Partitioning	Consistent hashing	Incremental scalability
High availability for writes	Vector clocks	(Check the paper)
Handling temporary failures	Sloppy quorum and hinted handoff	High availability and durability
Recovering from permanent failures	Merkle trees	(Check the paper)
Membership and failure detection	Gossip-based protocol	(Check the paper)

- All techniques are known in literature
- We focus on **consistent hashing** and **sloppy quorum**

Dynamo's Storage

➤ Problem:

- Key space of nodes = $\{0, 1, \dots, n-1\}$
- Set of objects $\{o_1, o_2, \dots, o_k\}$
- Create a hash function $h(x)$ such that
 - $h(x)$ maps an element in $\{o_1, o_2, \dots, o_k\}$ to number i
 - object x will be stored in node i

➤ Requirement of $h(x)$:

- Deterministic
- Efficient
- Evenly distributed among all nodes

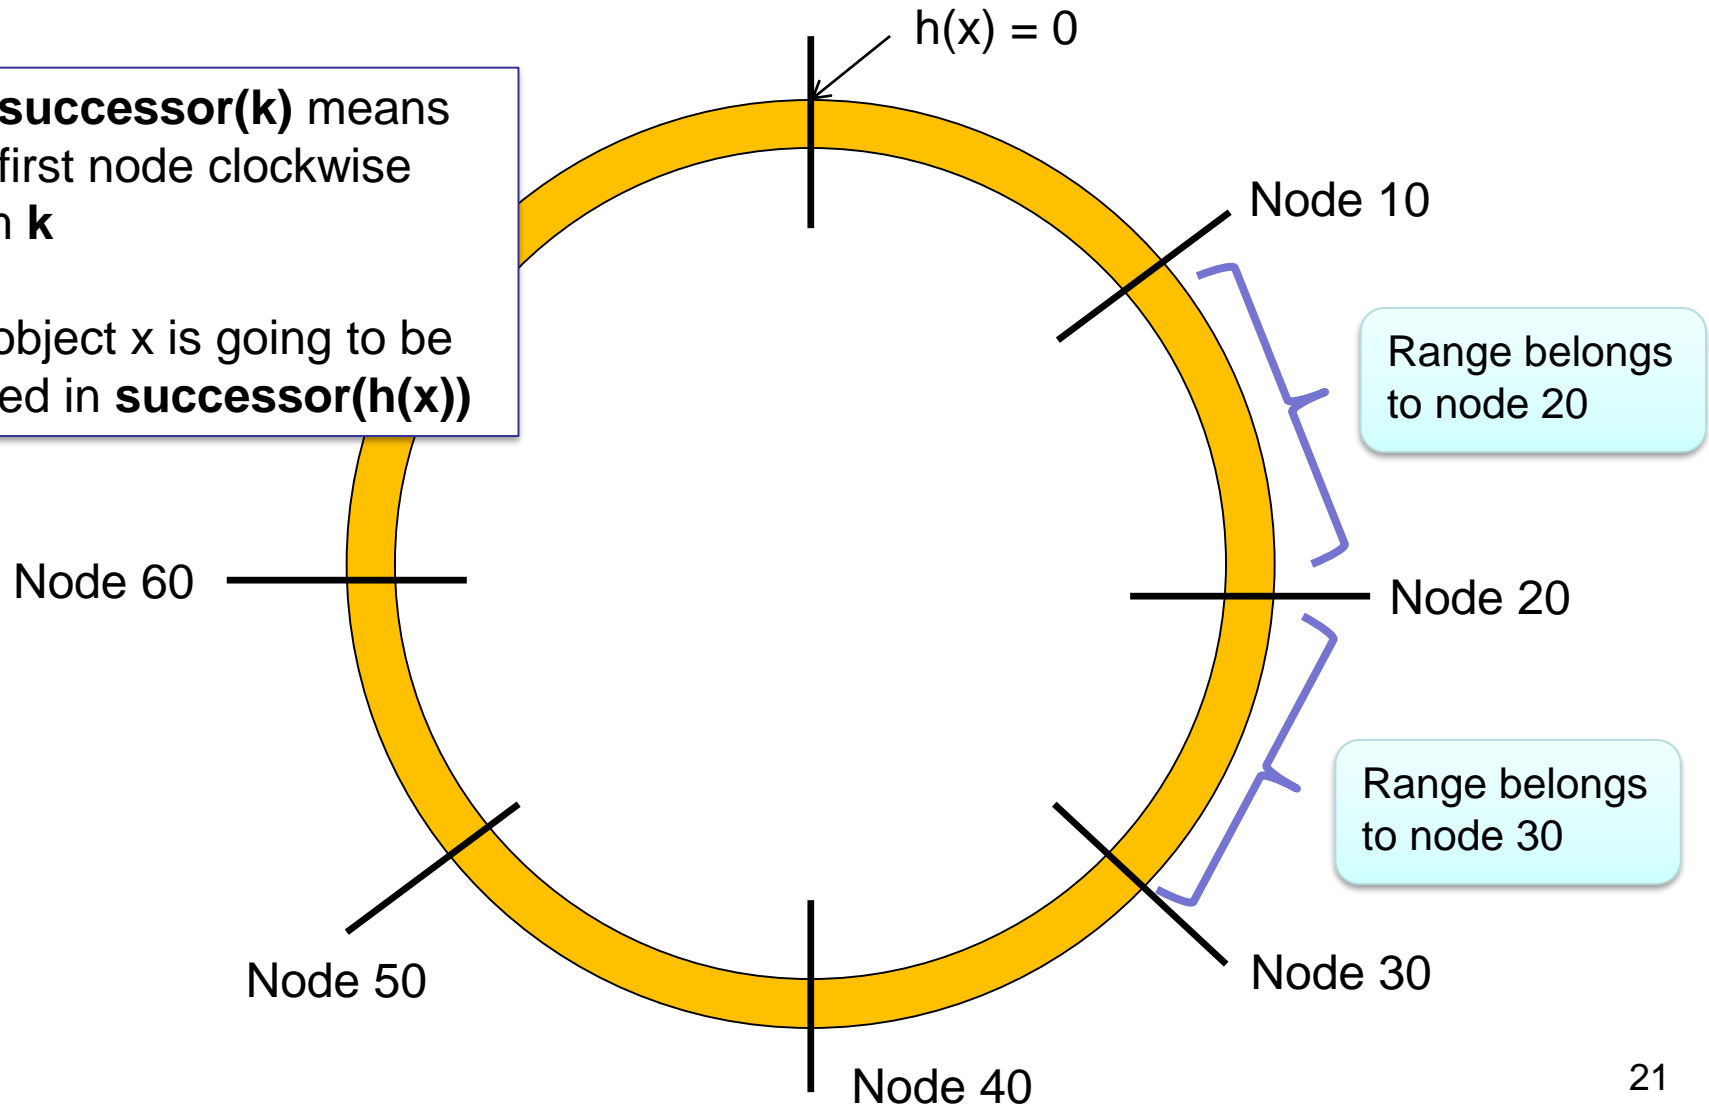
Dynamo's Storage

- In Dynamo, a key is hashed with **MD5** to create a 128-bit identifier, combined with modulo operation
- $h(x) = \text{MD5}(x) \bmod n$
- Any pros and cons?

Consistent Hashing: Store Objects

Let **successor(k)** means
the first node clockwise
from **k**

An object **x** is going to be
stored in **successor(h(x))**



Consistent Hashing: Node Join

- When a new server is added, only the data of the successor node is affected
- Example:
 - Node 26 is added
 - Data from node 30 is migrated to node 26
 - Update the routing tables (some routing protocol is needed)
- What about a node leave?

Consistent Hashing: Replication

- Replication: stores an object to N storage servers
 - 1st server is the one that is closest to $h(x)$ in the clockwise direction, i.e., $\text{successor}(h(x))$.
 - Then, $N-1$ more servers, who are the clockwise successors of $\text{successor}(h(x))$, will store a redundant copy of object x .
- Example: $h(x) = 15$ and $N = 3$. Who will store the objects?

Consistency

- Replication → consistency
- Challenge:
 - To decide which storage server(s) to store a particular object, without choosing all servers.
 - Choosing all servers implies **strict consistency**!
- Strict consistency:
 - Any read on a data item x returns a value corresponding to the result of the most recent write on x

Client-Side Consistency

➤ Strong consistency

- After the update reports a finished update, all subsequent access returns the updated value

➤ Eventual consistency

- After the update reports a finished update, the system may report the old value
- Nevertheless, the system will eventually report the updated value.

Server-Side Consistency

Configurable parameters

- **N** = the number of nodes that store replicas of the data
- **W** = the number of replicas that need to **acknowledge the receipt of update**
- **R** = the number of replicas that are **contacted when a data object is accessed**

Server-Side Consistency

- If $W + R > N$, then it is **strong consistency**
- If $N = 2$, $W = 2$, and $R = 1$, what will happen?
- If $N = 3$, $W = 2$, and $R = 2$, what may happen?
 - A write completes only when **any** two replica writes are confirmed
 - A read returns the values of **any** two replicas
 - Resolve conflict by data version

Versioning → Resolving Conflicts

- Not all updates may arrive at all replicas
- Application-based reconciliation
 - Each modification of data is treated as a new version
- Vector clocks are used for versioning
 - Capture causality between different versions of the same object
 - Vector clock is a set of (node, counter) pairs
 - Returned as a context from a get() operation

Server-Side Consistency

- If $W + R \leq N$, then it is **eventual consistency**
- If $N = 3$, $W = 2$, and $R = 1$, what will happen?
 - A high reading-demand configuration
 - Strong consistency is sacrificed

Dynamo's Setting

- Common (N,R,W) in Dynamo: $(3,2,2)$
- However, strict quorum is vulnerable to partitions
- **Sloppy quorum**
 - All reads/writes are performed on the first N **healthy** nodes from the preference list, which may not always be the first N nodes in the consistent hashing ring
 - e.g., if node A is down, a replica is sent to node A'
- **Hinted handoff**
 - Node A' receives a hint that the update is temporary
 - A' sends back the replica to A if A is recovered, and removes its replica

More on Eventual Consistency

➤ Causal consistency

- If process A has communicated to process B that it has updated a data item, a subsequent access by process B will return the updated value, and a write is guaranteed to supersede the earlier write

➤ Read-your-writes consistency

- Process A, after it has updated a data item, always accesses the updated value and will never see an older value
- A special case of the causal consistency model

More on Eventual Consistency

➤ Session consistency

- A practical version of previous model, where read-your-write consistency is guaranteed in the same session
- If a new session needs to be created, the guarantees do not overlap the sessions

➤ Monotonic read consistency

- If a process has seen a particular value for the object, any subsequent accesses will never return any previous values

➤ Monotonic write consistency

- Writes are serialized by the same process