

Lecture 12: Tail Tolerance

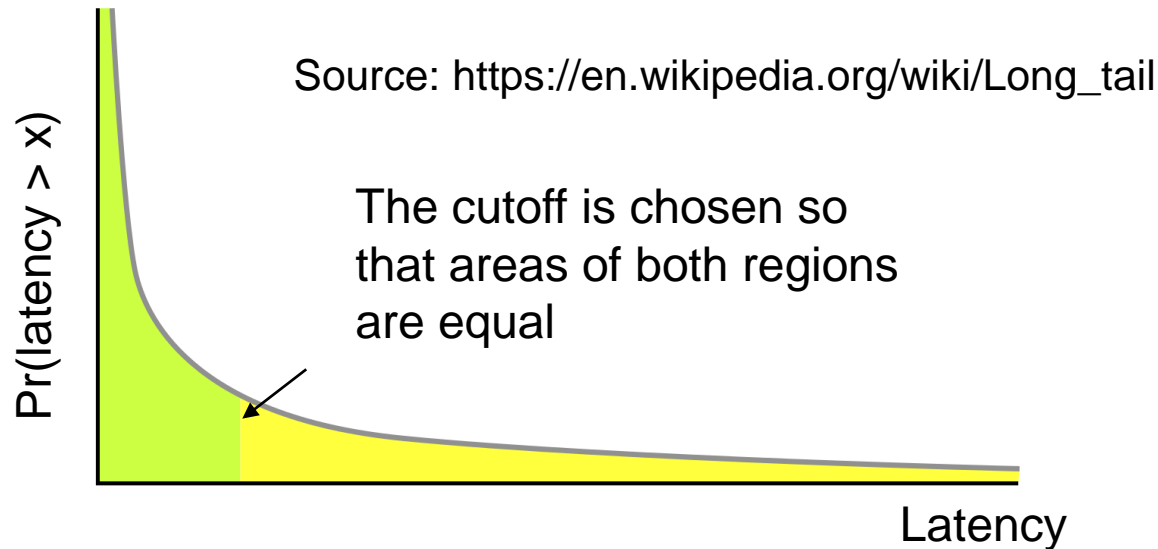
CSCI4180

Patrick P. C. Lee

Goal of this Lecture

- Why is tail latency bad?
- How Google handles tail latency?
- What are the tail latency issues in NetApp's data center storage?
- References:
 - Dean and Barroso, "The Tail at Scale", CACM 2013
 - Hao et al., "The Tail at Store: A Revelation from Millions of Hours of Disk and SSD Deployments", FAST 2016

What is a Tail?



- **Tail** refers to the latency in high percentiles
 - e.g., 95-percentile or 99-percentile latencies
- A tail is significant (or a heavy tail) if the high-percentile latency is high
 - Mathematically, $P(\text{latency} > x)$ is significant even for large x
- **Variability** of response time leads to high tail latency

Why Tails are Bad?

- Dominate the overall performance
 - e.g., stragglers in MapReduce determine the overall job performance
- Violation of service-level agreements
 - e.g., SLA may state 99.9% of requests are completed within 300 milliseconds
- Bad user experience
 - e.g., interactive queries

Why Tails are Bad?

- What should operators do with tails?
- Treat slow servers as hard failures?
 - Replace slow servers with new ones
 - Problem: there could be many false positives

Why Variability Exists?

- Resource sharing
 - CPUs / disks within machines
 - Switches or shared file systems
- Background jobs
 - Daemons
 - Maintenance
 - Garbage collection
- Queueing
- Energy management
 - Power limits of modern CPUs
 - Power saving modes

Variability Amplified at Scale

- Parallelization is often used to reduce latency, but also amplifies variability
- Scenario:
 - A request is broken into sub-operations that are issued to multiple servers in parallel
 - These sub-operations must all complete
 - Let $\Pr(\text{latency} > 1\text{s}) = 0.01$
 - Sub-operations of a request issued to 100 servers
 - $\Pr(\text{request} > 1\text{s}) = 1 - (1 - 0.01)^{100} = 0.63$
- What happen if a request is issued to one server?

Variability Amplified at Scale

	50%ile latency	95%ile latency	99%ile latency
One random leaf finishes	1ms	5ms	10ms
95% of all leaf requests finish	12ms	32ms	70ms
100% of all leaf requests finish	40ms	87ms	140ms

Measurements from a real Google service

- The slowest 5% of the requests to complete is responsible for half of the total 99%-percentile latency

Reducing Variability

- Differentiating service classes and higher-level queuing
- Reducing head-of-line blocking
 - HOL blocking: blocked by large requests at the front of a queue
 - Solution: break a long-running request into smaller slices and allow them interleave with others
- Managing background activities and synchronized disruption
 - Synchronizing all background activities reduces variability

Reducing Variability

- What about caching?
 - Cache the data that resides in slow nodes
- Caching is useful only when the entire working set can reside in a cache

Living with Variability

- Eliminating all latency variability is infeasible
 - Too large scale
 - Too complex
- Google develops tail-tolerant techniques that work around temporary latency hiccups

Within Request Short-Term Adaptations

- Replication of data items across servers
 - Provide both fault tolerance and tail tolerance
- **Hedged requests:**
 - Send a request to the primary server
 - Send the same request (i.e., hedged request) to a different server after some brief delay
 - Cancel any outstanding request if the response is received

Within Request Short-Term Adaptations

➤ Tied requests

- Enqueue copies of a request in multiple servers simultaneously
 - Allow servers to communicate updates on the status of these copies to each other (e.g., send a cancellation message when the request is processed)
- In Google's implementation of tied requests, medium latency is reduced by 16%, and 99.9-percentile latency is reduced by 40%

Cross-Request Long-Term Adaptations

➤ **Micro-partitions:**

- Generate many more partitions than servers
- Perform load balancing at the (more fine-grained) partition level
- Example: tablets in BigTable

➤ **Selective replication:**

- Detect or predict certain items (or micro-partitions) that are likely to cause load imbalance and create more replicas of them
- Spread the load of these “hot” items

Cross-Request Long-Term Adaptations

➤ Latency-induced probations:

- If a server is detected to respond slowly, put it on probation
- Exclude requests to the slow servers
- Issue shadow requests to the slow servers and collect statistics from them; bring the servers back to normal if they are ready

Special Cases

➤ Take “good enough” responses:

- Once a sufficient fraction of all servers responded, the user accepts the slightly incomplete (“good-enough”) results for better end-to-end latency
- Example: PageRank

➤ Canary requests:

- Rather than initially send a request to thousands of servers, send it first to one or two servers
- The remaining servers are only queried if there is a successful response from the canary in a reasonable period of time
- Rationale: don’t waste efforts

Summary

- Seven tail-tolerant techniques:
 - Reactive or proactive
- Rationale: optimize for the common case while providing resilience against uncommon cases

Tails in Storage


- What about the tail latency problem in storage environments?
- In reality, performance variability also exists in hard disks and solid-state drives (SSDs)
 - Hard disks
 - Weak disk head, bad packaging, missing screws, broken/old fans, too many disks/box, firmware bugs, bad sector remapping
 - Bandwidth drops by 80%, and introduces seconds of delay
 - SSDs:
 - Firmware bug, GC, ...
 - 4 – 100x slowdown

Field Study (NetApp)

- Study of over 450,000 disks and 4000 SSDs
 - Deployed as RAID groups
 - 87 days on average
 - Total: 857 million disk hours and 7 million SSD hours

	Disk	SSD
#RAID groups	38,029	572
#Data drives per group	3-26	3-22
#Data drives	458,482	4,069
Total drive hours	857,183,442	7,481,055
Total RAID hours	72,046,373	1,072,690

Metric

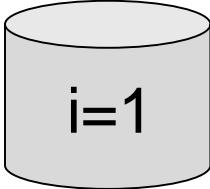
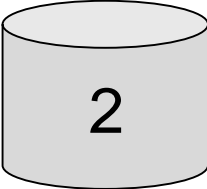
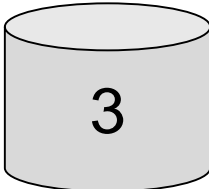
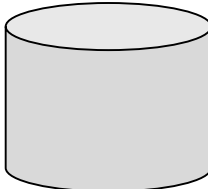
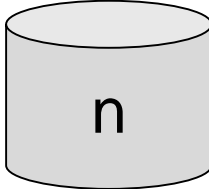


RAID group: P and Q drives are parity drives that provide redundancy.
Parities are **non-rotational**

Primary metric:

- L_i = average I/O latency per drive
 - L_{median} = median ($L_1 \dots L_N$)
- **Slowdown** $\rightarrow S_i = L_i / L_{\text{median}}$
 - **Slow drive** hour $\rightarrow \underline{S_i \geq 2x}$

Full-Stripe Workloads

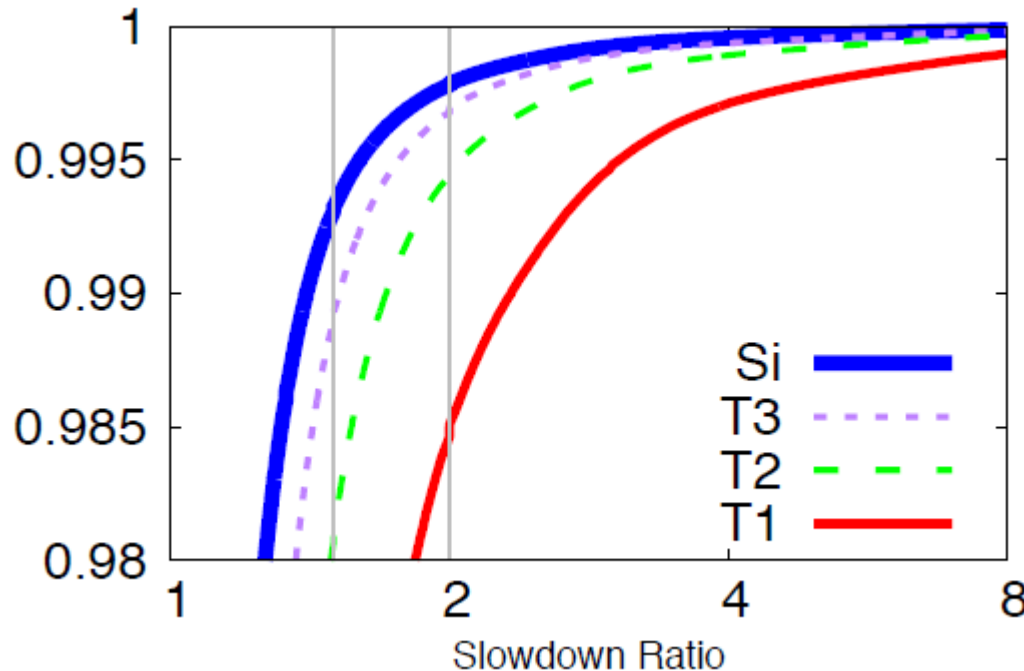
	RAID				
					
	i=1	2	3		n
Hourly average latency (ms)	10	9	22	8	11
Slowdown	1.0x	0.9x	2.2x	0.8x	1.1x
	L_{med}		T1		T2

- Most RAID I/O requests are **full-stripe**
- Three largest slowdowns (tails): T1, T2, T3
 - T1 = 2.2, T2 = 1.1, T3 = 1

Measurement Questions

- How often does slowdown happen?
- Does slowdown depend on the workloads?
- What are the possible root causes of the slowdown?
- What happens after slowdown occurs?

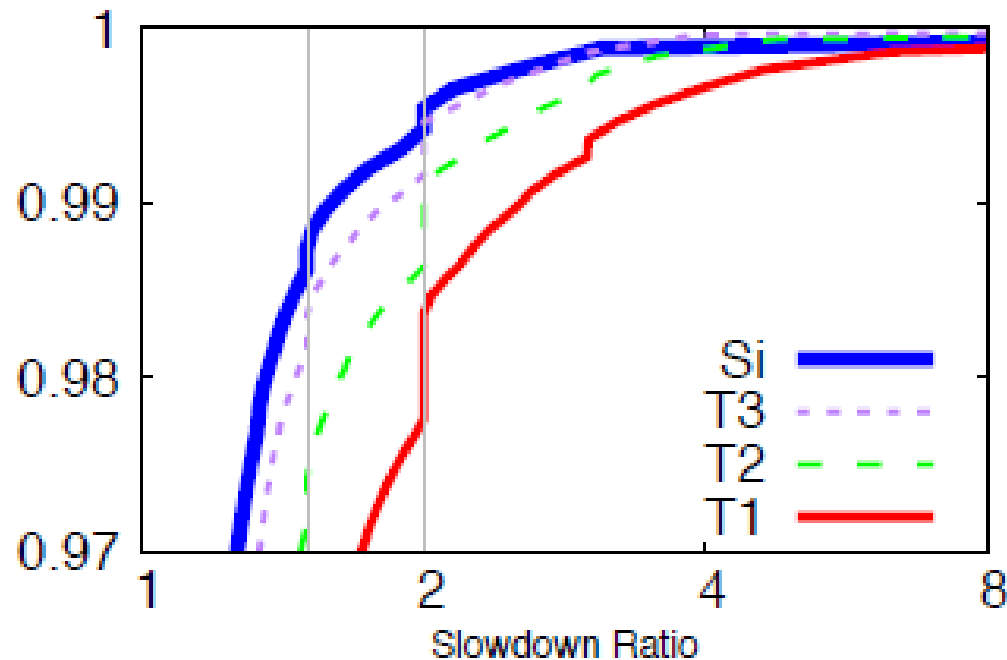
Slowdown Distribution (Disks)



Cumulative
distributions of
disk hours versus
slowdown

- Slowdowns are significant:
 - Some disks have $\geq 2x$ slowdown (i.e., $S_i \geq 2$) in 0.22% of time
 - In 1000 disk hours, 2 hours are slow
- Longest tails (T1) appear 1.5% of time

Slowdown Distribution (SSDs)



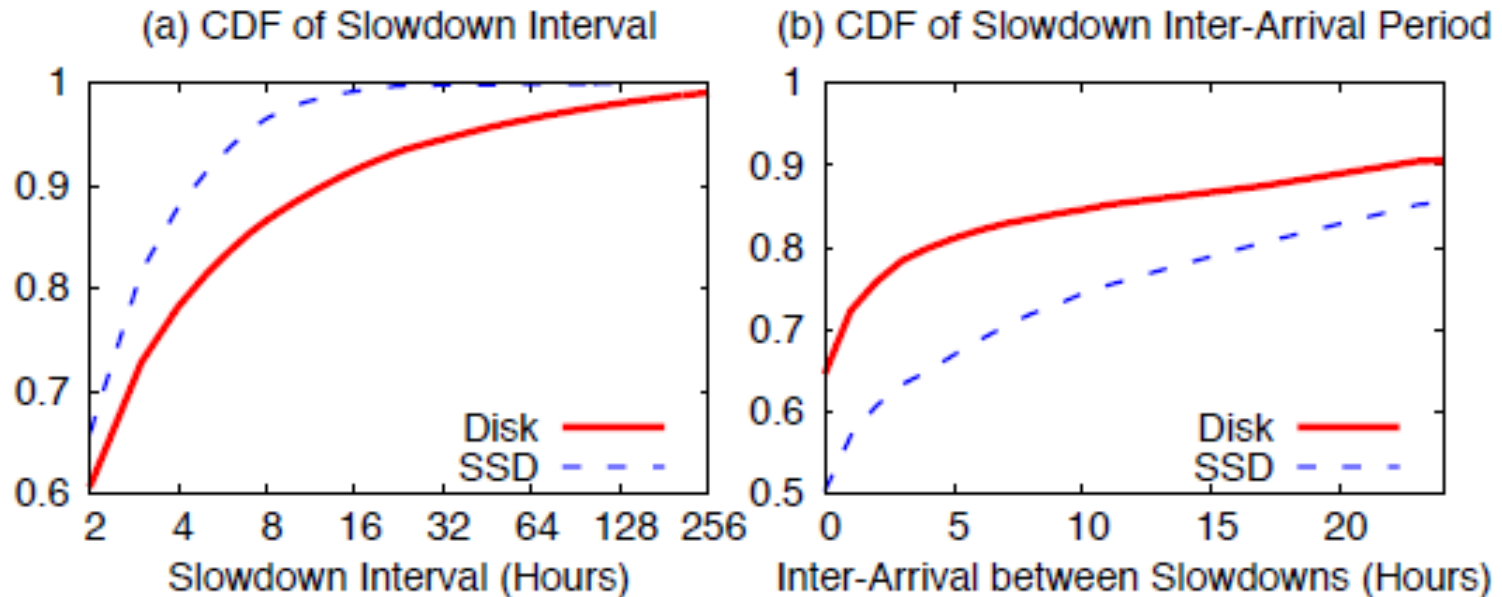
Cumulative
distributions of
disk hours versus
slowdown

- SSDs experience even more slowdown than disks (0.58% of time)
 - Faster devices don't imply fewer slowdowns
- Longest tails appear at 2.23% of time

Observations

- Storage performance instability is not uncommon
- Storage tails appear at a significant rate
- Tail-tolerant RAID has a significant potential to increase performance stability
 - Move the tail from T1 to T3

Temporal Behavior

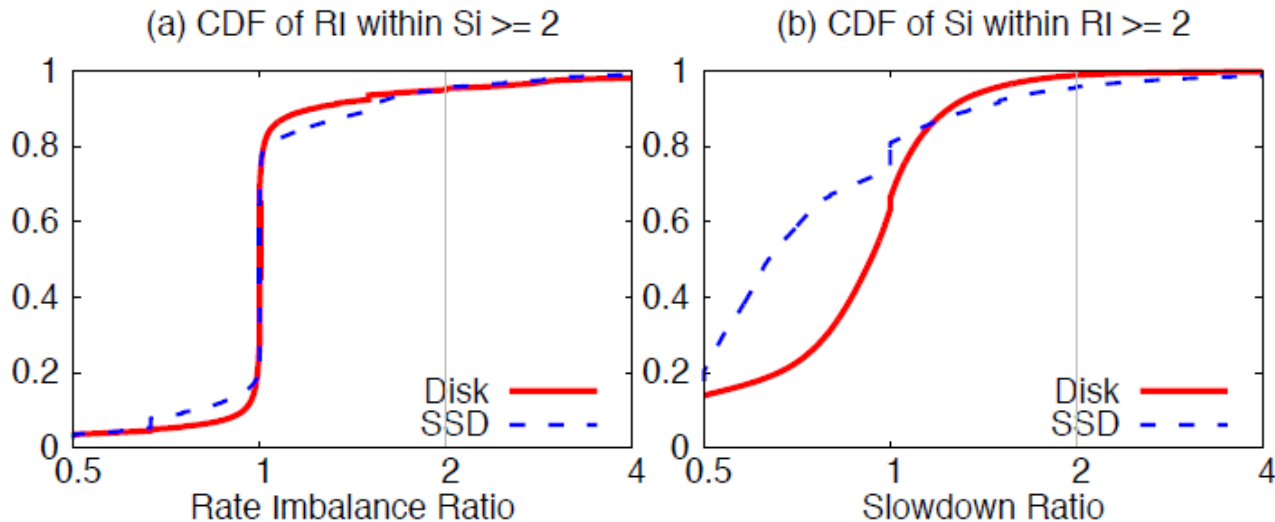


- Slowdown can persist over several hours
- Slowdown has a high temporal locality
 - 90% and 85% of disk and SSD slowdown occurrences from the same drive happen within the same day of the previous occurrence, respectively

Slowdown Population

- How many drives have experienced at least a slowdown in their lifetimes?
 - 25% of disks have seen $\geq 2x$ slowdown
 - 29% of SSDs have seen $\geq 2x$ slowdown
- A large proportion of drives experience slowdown

Workload Dependence



Slowdown vs I/O rate imbalance

- R_i = I/O rate of drive i
- Rate imbalance: $RI = \max(R_i) / R_{med}$
- (Left) Only 5% of slow drives have high RI
- (Right) Only 5% of imbalances occur in slow drives

Workload Dependence

- Similar observations for I/O size
- Conclusion: Slowdown is independent of I/O rate and size imbalance
- Further analysis shows that slowdown is a “silent” performance problem
 - Slowdowns are not accompanied with any explicit drive events

More Analysis

- Older disks tend to have more slowdowns, but no age correlation for SSDs
- Single-level-cell (SLC) SSDs slightly outperform multi-level-cell (MLC) SSDs in terms of performance stability
- Slowdown depends on SSD vendors
- No correlation with time of day
- Slow drive replacement rate is low
 - Unplug: 4-8% (within 24 hours after a slowdown)
 - Replug: 89-100% of unplugged drives are replugged

How to be Tail-Tolerant?

- **Reactive**: if a drive didn't return data after a timeout, perform an extra read to parity drives
- **Proactive**: perform extra reads to parity drives concurrently with original I/Os
- **Adaptive**: if the reactive policy is triggered repeatedly on the same drive, use the proactive policy

How to be Tail-Tolerant?

- Trade-offs:
 - Reactive: slow response time
 - Proactive: extra read traffic
- What about writes?
 - Can be mitigated by caching

Potential Directions

- More fine-grained analysis is needed
 - Traces are hourly averages
 - Absence of other metrics
- Correlations to failure events are important
- Tail mitigation techniques
 - What about distributed file systems (e.g., HDFS)?
- Can we *predict* tails?